



## Why probability?

- Probability is a foundational concept in inferential statistics.
- It is related to the idea that because it is hard to access the entire population, we can use a sample to make inferences about the population.
- We then take a sample, which is thought to be representative of the population, and use it to make inferences about the population.

# Why probability?

- However, even if we ensure the sample is similar to the population based on theoretical knowledge, we would never expect that it is exactly the same.
  - There will always be some difference, potentially due to random variation, which make the sample different from the population.
- Probability is the mathematical tool that allows us to quantify the uncertainty associated with the sample.
  - How far are we from the population?
  - If there are significant differences, are we looking for an unrepresentative sample?
  - Is a sample mean fundamentally different from the population mean?
- All the questions above can be answered using the concepts of probability.

## Basic probability concepts



# Approaches to probability calculation

- There are two main approaches to calculate probability:
  - Frequentist approach
  - Bayesian approach
- The frequentist approach is based on the idea that probability is the long-run frequency of an event.
  - For example, the probability of getting a head when flipping a coin is 0.5 because in the long run, half of the flips will be heads.
- This means that frequentist probability will calculate a probability based on past data

$$\text{Probability} = \frac{\text{Number of times an event occurred}}{\text{Total number of trials}}$$

# Approaches to probability calculation

- Another major approach to probability is the Bayesian approach (pronounced “bay-zee-an”).
- The Bayesian approach is based on the idea that probability is a measure of the degree of belief that an event will occur.
  - Hence, it is like a measure of how “good” our information is.
- We won’t focus Bayesian approaches on this module, but it is good to know that it exists, and researchers have devoted a lot of time to it.
  - Typically, we will want to understand frequentist approaches first and gain mathematical intuition before we move to Bayesian approaches.



# Basic probability rules

- Two major rules govern probability:
  - 1 All probabilities are between 0 and 1.
    - A probability of 0 means that the event will never happen (0% chance).
    - A probability of 1 means that the event will always happen (100% chance).
    - This means probabilities are always in proportion form, but can be easily converted to percentages.
  - 2 The sum of the probabilities for all possible outcomes of an event must sum to 1.
    - In an exam, you'd have two outcomes: pass or fail (assuming no other "weird" outcomes).
    - The probability of passing plus the probability of failing must sum to 1.
    - This means 100% of outcomes are covered in the "possible outcomes" set, each with a probability.





# Probability: the world of board games

- It is common to also see drawing a card from a deck of cards as an example of probability.
- Decks of cards have 52 cards, with 4 suits (hearts, diamonds, clubs, and spades) and 13 cards in each suit.
- Special cards include the King, Queen, the Jack, and the Ace.
  - There are 4 of each of these cards in the deck.
  - The probability of drawing a King is  $4/52$ , the probability of drawing a Queen is  $4/52$ , and so on.











# Union

- The union of two events A and B is the event that *either* A or B or both occur.
  - Meaning that any two of them can occur, but not necessarily both.
  - The union of A and B is denoted by  $A \cup B$  (set theory notation).
  - The probability of the union of A and B is denoted by  $P(A \cup B)$ .
- Example: what is the probability of getting one or two when rolling a die?
  - Notice the use of “or” in the question. Union is often associated with the “OR” logical operator.
  - This can be written as  $P(A \cup B)$ , where A is the event of getting a 1 and B is the event of getting a 2.



A Venn diagram illustrating the intersection of two sets,  $P$  and  $Q$ , within a universal set  $\xi$ .

- The universal set  $\xi$  is represented by the entire rectangle and contains the numbers: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16.
- Set  $P$  is represented by the orange circle and contains the numbers: 2, 4, 6, 8, 10, 12, 14, 16.
- Set  $Q$  is represented by the blue circle and contains the numbers: 5, 10, 15.
- The intersection of  $P$  and  $Q$  (the overlapping region) contains the number: 10.

Laboratorio de Investigación para el Desarrollo del Ecuador





# Complement

- The complement of an event  $A$  is the event that  $A$  does *not* occur.
  - Meaning that the event does not happen.
  - The complement of  $A$  is denoted by  $A'$  (set theory notation).
  - The probability of the complement of  $A$  is denoted by  $P(A')$ .
- Example: what is the probability of not getting a 1 when rolling a die?
  - The complement of getting a 1 is not getting a 1.
  - The complement of an event is often associated with the “NOT” logical operator.
- It is often very easy to calculate the complement of an event, as it is the *whole* minus the probability of the event.

$$P(A') = 1 - P(A)$$

# Complement in a Venn diagram




Figure 3: Complement in a Venn diagram. Source: BBC Bitesize





$$P(A \cup B) = P(A) + P(B)$$

where A and B are mutually exclusive events.

## Computing the probability of union: non-mutually exclusive events and the addition rule

- However, if the events are not mutually exclusive, we need to consider the probability of the intersection of the events.
- Events that might not be mutually exclusive happen very often. For instance, what is the probability of going to the beach and getting a sunburn?
  - These events are not mutually exclusive, as you can go to the beach and get a sunburn at the same time.

## Union under non-mutually exclusive events

- If we apply the “formula” above, we will overcount the probability of the intersection of the events.
- Example: the probability of finding a company that is both in Pichincha, and is listed as active.
  - If we know the probability of finding companies in Pichincha, it will include those listed as active and otherwise.
  - If we know the probability of finding companies listed as active, it will include those in Pichincha and otherwise.
- Summing these two without doing anything else will overcount the probability of finding a company that is both in Pichincha and is listed as active.
  - Solution: subtract the probability of the intersection of the events.

## Union under non-mutually exclusive events: the addition rule

- The probability of the union of two events is given by the addition rule:

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

- This is a broad formula that can be applied to any two events, whether they are mutually exclusive or not.
  - The intersection of two mutually exclusive events is 0, so the formula simplifies to the one we saw before.





- If two events are independent, the probability of the intersection of the events is given by the multiplication rule:

when  $A$  and  $B$  are independent events.





---

- For instance, we are interested in knowing the probability of finding an active company, *given* that the company is in Pichincha.
  - This is different from the probability of finding an active company in general.
- This is a conditional probability: the condition is that the company is in Pichincha.
  - This is denoted by  $P(A|B)$ , where A is the event of finding an active company and B is the event of finding a company in Pichincha.



## Conditional probability example

- Following our above example:
  - A is the event of finding an active company.
  - B is the event of finding a company in Pichincha.
- We find active companies in Pichincha (the intersection of the events) and divide by the probability of finding a company in Pichincha.
- Conditionality is the theoretical probability concept of “subsetting” a sample space based on the occurrence of another event.
  - So, we are theorizing our use of `filter()` in R.





## Laboratorio de Investigación para el Desarrollo del Ecuador



# Random variables and probability distributions











---





- The fair die example (rolling a die) has a more complex PDF. See below:

Value of the random variable	Probability
1	$1/6$
2	$1/6$
3	$1/6$
4	$1/6$
5	$1/6$
6	$1/6$

- In this case, the value of the random variable  $X$  is the number that comes up when rolling a die.
- The probability of getting a 1 is  $1/6$ , the probability of getting a 2 is









## Example: expected value of a coin

- For the coin example, the expected value is:

$$E(X) = 0 \cdot 0.5 + 1 \cdot 0.5 = 0.5$$

- This means that the expected value of the number of heads when flipping a coin is 0.5.
- Does this make sense? It will often happen that the expected value is not a possible value of the random variable.
  - This is because the expected value is a measure of central tendency, and not necessarily a possible value of the random variable.
- In an intuitive sense, the expected value is the value we would expect to get if we repeated the experiment many, many times, and took the average of the outcomes.



## Example: variance of a coin

- For the coin example, the variance is:

$$\text{Var}(X) = (0 - 0.5)^2 \cdot 0.5 + (1 - 0.5)^2 \cdot 0.5 = 0.25$$

- This means that the variance of the number of heads when flipping a coin is 0.25.
- Once again, higher variance means that the values of the random variable are more spread out.

## Laws of expected value and variance

- When working with expected values and variances, there are some laws that can be useful as shorthands.
- For any constants  $a$  and  $b$ :

$$E(aX + b) = aE(X) + b$$

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

# Continuous random variables

- There is a specific detail about the PDF of continuous random variables: the probability of the random variable taking on a specific value is 0.
- This is because continuous random variables can take on an infinite number of values, and the probability of any one value is 0.
  - Why infinite? Because we can always find a value between two values (i.e. with a lot of decimals).
- Intuitively, think that whenever we have a continuous random variable, it is extremely unlikely that we will get a specific value.
  - Measurements are never perfect. We can never measure a person's height to the exact millimeter, hence, we can't even know the exact value of the random variable.





## The CDF of continuous RVs

- The CDF of a continuous random variable is a function that describes the probability that the random variable is less than or equal to a certain value.
- In this context, this can be understood as a function  $F(x)$  that always gives you the area under the curve of the PDF up to a certain value  $x$ .
- Hence, the CDF is an integral of the PDF.

$$F(x) = \int_{-\infty}^x f(t)dt$$

where  $f(t)$  is the PDF of the random variable.

## Expected value of a continuous RV

- The expected value of a continuous random variable is calculated in the same way as for discrete random variables.
- We multiply each possible value of the random variable by its probability, and sum these products. However, this time, we use an integral instead of a sum.

$$E(X) = \int_{-\infty}^{\infty} x \cdot f(x) dx$$

- Don't worry about these! We need not calculate these by hand. In an applied context, we just need to know what PDFs and CDFs are.



- There are several continuous probability distributions, but the most common is the normal distribution
  - Also known as the Gaussian distribution or the bell curve.
- **Why is the normal distribution the most common?**
  - 1 Flexibility and simplicity:**
    - The normal distribution is easy to handle and does not require complex mathematical tools.
  - 2 Common in real-world data:**
    - Many variables are approximately normally distributed, making it practical for various applications.

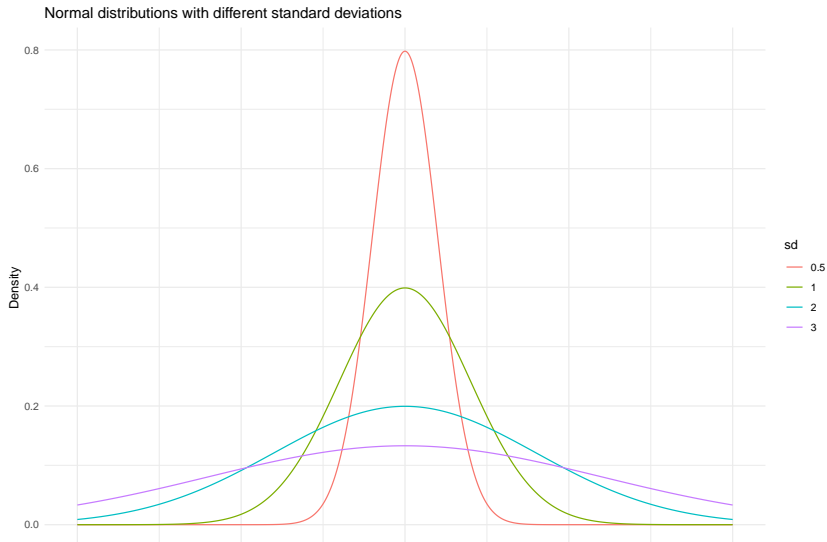




- Different values of  $\sigma$  affect how flat or peaked the distribution is.
  - A small  $\sigma$  means the distribution is tall and thin.
  - A large  $\sigma$  means the distribution is short and wide.
- Larger values of  $\sigma$  mean that the distribution is more spread out.
  - Less data will be close to the mean, and more data will be far from the mean.



# Different normal distributions under varying $\sigma$










# The CDF of the normal distribution

- The CDF of the normal distribution is given by the formula:

$$F(x) = \int_{-\infty}^x f(t)dt$$

- We never need to do this by hand! We can use software tools to calculate the CDF of the normal distribution.
- However, it must be clear that CDFs are often defined as the area under the curve up to a certain value of the random variable.
  - They answer the question “what is the probability that the random variable is less than or equal to a certain value?”

Normal Distribution (mean = 0, sd = 1)




- The empirical or the 68-95-99.7 rule is a rule of thumb that describes the proportion of values that fall within a certain number of standard deviations from the mean in a normal distribution.
- It is a direct consequence of the properties of the normal distribution, and the use of the CDF for probabilities which are easier to calculate.





# The empirical rule



# The standard normal distribution

- The **standard normal distribution** is a normal distribution with  $\mu = 0$  and  $\sigma = 1$ .
  - This is a very common distribution in statistics.
- The PDF of the standard normal distribution is given by the formula:

$$f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}$$

- This is the normal distribution with  $\mu = 0$  and  $\sigma = 1$ .
- The standard normal distribution is often denoted by  $Z$ .

## Standardizing normal distributions

- All normal distributions can be standardized to the standard normal distribution.
- We do this with our beloved  $z$ -scores.

$$z = \frac{x - \mu}{\sigma}$$

- This  $z$ -score tells us how many standard deviations a value is from the mean.
- So, if a distribution is normal and we transform it to a standard normal distribution, we can use the standard normal distribution to calculate probabilities.

- R has built-in functions to work with the normal distribution.
- The `dnorm()` function calculates the PDF of the normal distribution.
  - It takes the value of the random variable, the mean, and the standard deviation as arguments.
  - The use of `d` in the function name is a convention in R to denote the density function.
- We would use `dnorm()` to calculate the likelihood of different outcomes of the random variable.
  - Also to plot the PDF of the normal distribution.



pnorm()

- **Exercise 1:** Calculate the probability of a standard normal random variable being less than 1.96.
- **Exercise 2:** Calculate the probability of a standard normal random variable being greater than -1.64.
- **Exercise 3:** Calculate the probability of a standard normal random variable being between -1.96 and 1.96.



## qnorm()

- Consider the following for a std. normal RV:
- **Exercise 1:** Find the value of a standard normal random variable that corresponds to the 95th percentile.
- **Exercise 2:** Find the value of a standard normal random variable that corresponds to the 10th percentile.
- **Exercise 3:** What RV value corresponds to a cumulative probability of 0.5?



