

Why probability?

- Probability is a foundational concept in inferential statistics.
- It is related to the idea that because it is hard to access the entire population, we can use a sample to make inferences about the population.
- We then take a sample, which is thought to be representative of the population, and use it to make inferences about the population.

Why probability?

- However, even if we ensure the sample is similar to the population based on theoretical knowledge, we would never expect that it is exactly the same.
 - There will always be some difference, potentially due to random variation, which make the sample different from the population.
- Probability is the mathematical tool that allows us to quantify the uncertainty associated with the sample.
 - How far are we from the population?
 - If there are significant differences, are we looking for an unrepresentative sample?
 - Is a sample mean fundamentally different from the population mean?
- All the questions above can be answered using the concepts of probability.

Basic probability concepts

Basic probability rules

- Two major rules govern probability:
- 1 All probabilities are between 0 and 1.
 - A probability of 0 means that the event will never happen (0% chance).
 - A probability of 1 means that the event will always happen (100% chance).
 - This means probabilities are always in proportion form, but can be easily converted to percentages.
 - 2 The sum of the probabilities for all possible outcomes of an event must sum to 1.
 - In an exam, you'd have two outcomes: pass or fail (assuming no other "weird" outcomes).
 - The probability of passing plus the probability of failing must sum to 1.
 - This means 100% of outcomes are covered in the "possible outcomes" set, each with a probability.

- The other common example is rolling a die.
- A die has 6 faces, each with a number from 1 to 6.
- The probability of getting a 1 is $1/6$, the probability of getting a 2 is $1/6$, and so on.

The lingo: experiments, events, outcomes and sample spaces

- In simple experiments, we often see “events” covering one outcome, and not more.
 - However, events can “group” outcomes in a specific way.
 - E.g. the event of getting an even number when rolling a die. This event covers outcomes 2, 4, and 6.
- The **sample space** is the set of all possible outcomes of an experiment.
 - For example, when flipping a coin, the sample space is $\{H, T\}$.
 - When rolling a die, the sample space is $\{1, 2, 3, 4, 5, 6\}$.

Set theory and probability

- Probability is closely related to set theory.
 - Sets are collections of objects, and in probability, we often think of events as sets of outcomes.
- Thus, it is common to see notation from set theory in probability.
 - A set is denoted by curly braces, e.g. $\{1, 2, 3, 4, 5, 6\}$.
 - We use letters to denote sets, e.g. A, B, C, etc.
 - We use set operations to calculate probabilities.
 - We use Venn diagrams to visualize probabilities.
- Set theory operators, such as belongs to (\in), union (\cup), intersection (\cap), and complement ($'$) are used in probability.
 - These often also come in terms of logical operators, such as “and” and “or”.

A Venn diagram illustrating the intersection of two sets, P and Q , within a universal set ξ .


- Set P (orange circle) contains the elements: 2, 4, 6, 8, 12, 14, 16.
- Set Q (blue circle) contains the elements: 5, 15.
- The intersection of P and Q contains the element: 10.
- The universal set ξ (the entire box) contains all elements: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16.

Laboratorio de Investigación para el Desarrollo del Ecuador

Complement

- The complement of an event A is the event that A does *not* occur.
 - Meaning that the event does not happen.
 - The complement of A is denoted by A' (set theory notation).
 - The probability of the complement of A is denoted by $P(A')$.
- Example: what is the probability of not getting a 1 when rolling a die?
 - The complement of getting a 1 is not getting a 1.
 - The complement of an event is often associated with the “NOT” logical operator.
- It is often very easy to calculate the complement of an event, as it is the *whole* minus the probability of the event.

$$P(A') = 1 - P(A)$$



$$P(A \cup B) = P(A) + P(B)$$

where A and B are mutually exclusive events.

Computing the probability of union: non-mutually exclusive events and the addition rule

- However, if the events are not mutually exclusive, we need to consider the probability of the intersection of the events.
- Events that might not be mutually exclusive happen very often. For instance, what is the probability of going to the beach and getting a sunburn?
 - These events are not mutually exclusive, as you can go to the beach and get a sunburn at the same time.

Union under non-mutually exclusive events: the addition rule

- The probability of the union of two events is given by the addition rule:

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

- This is a broad formula that can be applied to any two events, whether they are mutually exclusive or not.
 - The intersection of two mutually exclusive events is 0, so the formula simplifies to the one we saw before.

Independence

- Two events are independent if the occurrence of one event does not affect the occurrence of the other event.
- For example, the probability of getting a head when flipping a coin is independent of the probability of getting a 1 when rolling a die.
 - Basically means the events are not related.
- If two events are independent, the probability of the intersection of the events is the product of the probabilities of the events.
 - This is called the multiplication rule.

- If two events are independent, the probability of the intersection of the events is given by the multiplication rule:

$$P(A \cap B) = P(A) \times P(B)$$

when A and B are independent events.

Conditional probability

- For instance, we are interested in knowing the probability of finding an active company, *given* that the company is in Pichincha.
 - This is different from the probability of finding an active company in general.
- This is a conditional probability: the condition is that the company is in Pichincha.
 - This is denoted by $P(A|B)$, where A is the event of finding an active company and B is the event of finding a company in Pichincha.

Conditional probability example

- Following our above example:
 - A is the event of finding an active company.
 - B is the event of finding a company in Pichincha.
- We find active companies in Pichincha (the intersection of the events) and divide by the probability of finding a company in Pichincha.
- Conditionality is the theoretical probability concept of “subsetting” a sample space based on the occurrence of another event.
 - So, we are theorizing our use of `filter()` in R.

- Conditional probability is a key concept for independence: if two events are truly independent, the conditional probability of one event given the other is the same as the probability of the event (without the condition).
 - This is since conditioning on the other event does not change the probability of the event.
- Example: if we know that a die roll gives an even number, what is the probability of getting a 2 in a second roll?
 - Two die rolls are independent events (from our logical understanding of the experiment)
 - So, the probability of getting a 2 in the second roll is the same as the probability of getting a 2 any time we roll a die.

Laboratorio de Investigación para el Desarrollo del Ecuador

The multiplication rule for independent events

- For instance, if we know the probability of getting a head when flipping a coin is 0.5, and the probability of getting a 1 when rolling a die is $1/6$, we can calculate the probability of getting a head and a 1.

$$P(\text{Head} \cap 1) = P(\text{Head}) \times P(1) = 0.5 \times \frac{1}{6} = \frac{1}{12}$$

- This is a very powerful rule that can be used to calculate the probability of the intersection of independent events.
- Notice that I don't use my companies' example here, as the events are likely not independent.
 - Active companies may tend to be in Pichincha as it is the capital
 - This underscores the need to actually understand the nature of the events, or the *data generating process*.

Random variables and probability distributions

- The fair die example (rolling a die) has a more complex PDF. See below:

Value of the random variable	Probability
1	$1/6$
2	$1/6$
3	$1/6$
4	$1/6$
5	$1/6$
6	$1/6$

- In this case, the value of the random variable X is the number that comes up when rolling a die.
- The probability of getting a 1 is $1/6$, the probability of getting a 2 is

- For the coin example, the expected value is:

$$E(X) = 0 \cdot 0.5 + 1 \cdot 0.5 = 0.5$$

- This means that the expected value of the number of heads when flipping a coin is 0.5.
- Does this make sense? It will often happen that the expected value is not a possible value of the random variable.
 - This is because the expected value is a measure of central tendency, and not necessarily a possible value of the random variable.
- In an intuitive sense, the expected value is the value we would expect to get if we repeated the experiment many, many times, and took the average of the outcomes.

Example: variance of a coin

- For the coin example, the variance is:

$$\text{Var}(X) = (0 - 0.5)^2 \cdot 0.5 + (1 - 0.5)^2 \cdot 0.5 = 0.25$$

- This means that the variance of the number of heads when flipping a coin is 0.25.
- Once again, higher variance means that the values of the random variable are more spread out.

Laws of expected value and variance

- When working with expected values and variances, there are some laws that can be useful as shorthands.
- For any constants a and b :

$$E(aX + b) = aE(X) + b$$

$$\text{Var}(aX + b) = a^2 \text{Var}(X)$$

The CDF of continuous RVs

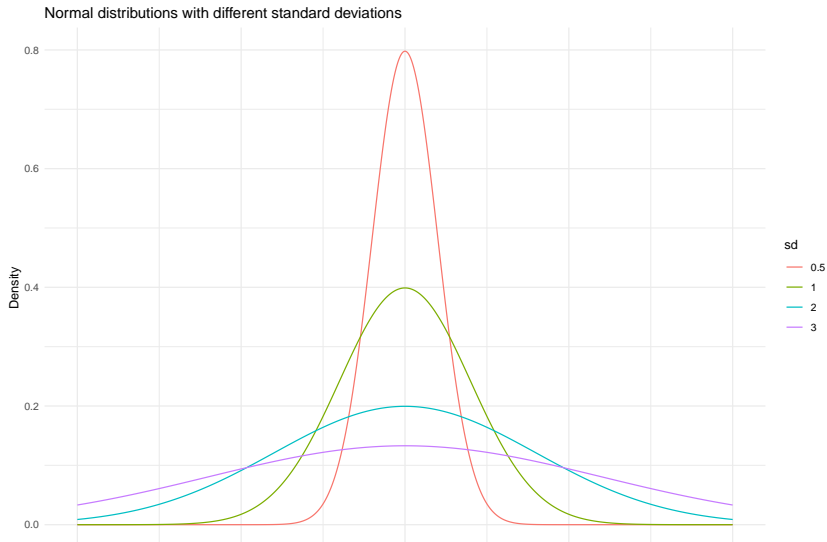
- The CDF of a continuous random variable is a function that describes the probability that the random variable is less than or equal to a certain value.
- In this context, this can be understood as a function $F(x)$ that always gives you the area under the curve of the PDF up to a certain value x .
- Hence, the CDF is an integral of the PDF.

$$F(x) = \int_{-\infty}^x f(t)dt$$

where $f(t)$ is the PDF of the random variable.


- There are several continuous probability distributions, but the most common is the normal distribution
 - Also known as the Gaussian distribution or the bell curve.
- **Why is the normal distribution the most common?**
 - 1 **Flexibility and simplicity:**
 - The normal distribution is easy to handle and does not require complex mathematical tools.
 - 2 **Common in real-world data:**
 - Many variables are approximately normally distributed, making it practical for various applications.

Different normal distributions under varying σ






Normal Distribution (mean = 0, sd = 1)



- The empirical or the 68-95-99.7 rule is a rule of thumb that describes the proportion of values that fall within a certain number of standard deviations from the mean in a normal distribution.
- It is a direct consequence of the properties of the normal distribution, and the use of the CDF for probabilities which are easier to calculate.

The empirical rule



Using R with the normal distribution

- R has built-in functions to work with the normal distribution.
- The `dnorm()` function calculates the PDF of the normal distribution.
 - It takes the value of the random variable, the mean, and the standard deviation as arguments.
 - The use of `d` in the function name is a convention in R to denote the density function.
- We would use `dnorm()` to calculate the likelihood of different outcomes of the random variable.
 - Also to plot the PDF of the normal distribution.

pnorm()

- **Exercise 1:** Calculate the probability of a standard normal random variable being less than 1.96.
- **Exercise 2:** Calculate the probability of a standard normal random variable being greater than -1.64.
- **Exercise 3:** Calculate the probability of a standard normal random variable being between -1.96 and 1.96.

- The `qnorm()` function calculates the quantiles of the normal distribution.
 - It takes the probability, the mean, and the standard deviation as arguments.
 - The use of `q` in the function name is a convention in R to denote the quantile function.
- This is useful when we want to find the value of the random variable that corresponds to a certain probability.
- This is otherwise known as the inverse CDF

Using R with the normal distribution: exercises with `qnorm()`

- Consider the following for a std. normal RV:
- **Exercise 1:** Find the value of a standard normal random variable that corresponds to the 95th percentile.
- **Exercise 2:** Find the value of a standard normal random variable that corresponds to the 10th percentile.
- **Exercise 3:** What RV value corresponds to a cumulative probability of 0.5?

